

POSIX threads (pthreads)

POSIX threads: pthreads

- POSIX: Portable Operating System Interface for UNIX – Interface to Operating System utilities
- Pthreads: The POSIX threading interface
 - Thread implementations adhering to the POSIX standard
 - System calls to create and synchronize threads
 - Should be relatively uniform across UNIX-like OS platforms
- Pthreads contain support for
 - Creating parallelism
 - Synchronization
- No explicit support for communication, because shared memory is implicit

Creating and terminating pthreads

`pthread_create(thread, attribute, thread_function, arg)`

- `thread`: unique identifier for the new thread (thread id or handle)
- `attribute`: object that may be used to specify attribute
 - example attribute: minimum stack size
- `thread_function`: function to be executed when thread is created
- `arg`: argument to be passed to `thread_function` when it starts
- returns `error_code`: set to nonzero if the create operation fails

`pthread_exit(status), pthread_attr_init(attr),
pthread_attr_destroy(attr)`

Example: creating and terminating pthreads

```
#include <pthread.h>
#include <stdio.h>
#define NUM_THREADS      4

void *PrintHello(void *threadid)
{
    long tid;
    tid = (long)threadid;
    printf("Hello World! It's me, thread #%ld!\n", tid);
    pthread_exit(NULL);
}

int main (int argc, char *argv[])
{
    pthread_t  threads[NUM_THREADS];
    int rc;
    long t;
    for(t=0; t<NUM_THREADS; t++){
        printf("In main: creating thread %ld\n", t);
        rc = pthread_create(&threads[t], NULL, PrintHello, (void *)t);
        if (rc){
            printf("ERROR; return code from pthread_create() is %d\n", rc);
            exit(-1);
        }
    }
    /* Last thing that main() should do */
    pthread_exit(NULL); /* main will block and wait until all threads are done */
}
```

Example: thread argument passing

```
struct thread_data{
    int  thread_id;
    int  sum;
    char *message;
};

struct thread_data thread_data_array[NUM_THREADS];

void *PrintHello(void *threadarg)
{
    struct thread_data *my_data;
    ...
    my_data = (struct thread_data *) threadarg;
    taskid = my_data->thread_id;
    sum = my_data->sum;
    hello_msg = my_data->message;
    ...
}

int main (int argc, char *argv[])
{
    ...
    thread_data_array[t].thread_id = t;
    thread_data_array[t].sum = sum;
    thread_data_array[t].message = messages[t];
    rc = pthread_create(&threads[t], NULL, PrintHello, (void *) &thread_data_array[t]);
    ...
}
```

- *each thread receives a unique instance of the data structure*

Thread management

- `pthread_join(thread_id, status)`
 - blocks the calling thread until the specified `thread_id` terminates
- `pthread_yield()`
 - Informs the scheduler that the thread is willing to yield its quantum, requires no argument
- `pthread_detach(thread);`
 - Informs the library that the thread's exit status will not be needed by subsequent `pthread_join` calls, resulting in better thread performance.

MASTER
THREAD

`pthread_create()`

`pthread_join()`

WORKER
THREAD

`pthread_exit()`



pthread attributes

- Once an initialized object attribute exists, changes can be made.
- To change the stack size of a thread to 8192 (before calling `pthread_create`):
 - `pthread_attr_setstacksize(&my_attributes, (size_t)8192);`
- To get the stack size:
 - `size_t my_stack_size;
pthread_attr_getstacksize(&my_attributes, &my_stack_size);`

Other pthread attributes

- Detached state: set if no other thread will use pthread_join to wait for this thread (improves efficiency)
- Guard size: use to protect against stack overflow
- Inherit scheduling attributes (from creating thread) – or not
- Scheduling parameter(s) – in particular, thread priority
- Scheduling policy: FIFO or Round Robin
- Contention scope: with what threads does this thread compete for a CPU
- Stack address: explicitly dictate where the stack is located
- Lazy stack allocation: allocate on demand (lazy) or all at once, “up front”

Shared data

- Variables declared outside of main are shared
- Objects allocated on the heap may be shared (if pointer is passed)
- Variables on the stack are private: passing pointer to these to other threads can cause problems

Loop-level parallelism

- Many applications have parallelism in loops

```
double stuff [n][n];  
for (int i = 0; i < n; i++)  
    for (int j = 0; j < n; j++)  
        ...  
        pthread_create (... , update_stuff, ..., &stuff[i][j]);
```

- But overhead of thread creation is nontrivial
 - update_stuff should have a significant amount of work
- Common performance pitfall: too many threads
 - Cost of creating a thread is 10s of thousands of cycles on modern architectures
 - Solution: “threadblocking”: use a small # of threads, often equal to the number of cores/processors or hardware threads
 - “block” of loop iterations executed by each thread

Example of parallel execution: “blocked” threads

| thread 0 | thread 1 | thread 2 | thread 3 |
|-------------------|---------------------|---------------------|---------------------|
| $i = 0 \dots 127$ | $i = 128 \dots 255$ | $i = 256 \dots 383$ | $i = 384 \dots 511$ |
| $A[i]$ | $A[i]$ | $A[i]$ | $A[i]$ |
| + | + | + | + |
| $B[i]$ | $B[i]$ | $B[i]$ | $B[i]$ |
| = | = | = | = |
| $C[i]$ | $C[i]$ | $C[i]$ | $C[i]$ |

Data races

- Common correctness pitfall: race conditions
- A race condition or data race occurs when:
 - Two threads access the same variable, and at least one does a write
 - The accesses are concurrent (not synchronized) so they could happen simultaneously

```
static int s = 0;
```

```
// Thread 1  
for (i=0; i<n/2; i++)  
    s = s + 1
```

```
// Thread 2  
for (i=n/2; i<n; i++)  
    s = s + 1
```

Example of data race

- Increments may be lost if both threads read simultaneously
- Need to make the read+increment+write an “atomic” operation
- Solution: locks aka Mutexes

```
static int s = 0;
```

```
// Thread 1 (core 1)  
read s into R1  
increment R1  
write R1 to s
```

```
// Thread 2 (core 2)  
read s into R2  
increment R2  
write R2 to s
```

Basic synchronization: mutexes

- Mutexes -- mutual exclusion aka locks
 - threads are working mostly independently
 - need to access common data structure

```
lock_t *lock = alloc_and_init();  
acquire(lock);  
access data  
release(lock);
```

Mutexes in pthreads

- To create a mutex:

```
#include <pthread.h>
pthread_mutex_t amutex = PTHREAD_MUTEX_INITIALIZER;
pthread_mutex_init(&amutex, NULL);
```

- To use it:

```
int pthread_mutex_lock(amutex);
int pthread_mutex_unlock(amutex);
```

- To deallocate it:

```
int pthread_mutex_destroy(pthread_mutex_t *mutex);
```

Correctness pitfalls with locks

- Locks cover too small a region

```
acquire (a)
tmp = s
release (a)
tmp++           /* updates can still be lost */
acquire (a)
s = tmp
release (a)
```

- Multiple locks may be held, but can lead to deadlock:

| | |
|----------|----------|
| thread1 | thread2 |
| lock(a); | lock(b); |
| lock(b); | lock(a); |

Performance pitfalls with locks

- Critical region is too large
 - Little or no true parallelism
 - Lock cost can go up with more contention
 - Solution: make critical regions as small as possible (but no smaller)
 - Solution: Use different locks for different data structures
- Locks themselves may be expensive
 - The overhead of locking / unlocking can be high

Basic synchronization mechanisms: barriers

- Barrier -- global synchronization
- Especially common when running multiple copies of the same function in parallel (SPMD)
- simple use of barriers -- all threads hit the same one

```
work_on_my_subgrid();  
barrier;  
read_neighboring_values();  
barrier;
```

- more complicated -- barriers on branches (or loops)

```
if (thread_id % 2 == 0) {  
    work1();  
    barrier;  
} else {  
    barrier;  
}
```

- barriers are not provided in all thread libraries

Creating and initializing a barrier

- To (dynamically) initialize a barrier, use code similar to this (sets the number of threads to 4):

```
pthread_barrier_t b;  
pthread_barrier_init(&b, NULL, 4);
```

- The second argument specifies an object attribute; using NULL yields the default attributes.
- other attributes: PTHREAD_PROCESSOR_SHARED/
PTHREAD_PROCESSOR_PRIVATE
- To wait at a barrier, a process executes:

```
pthread_barrier_wait(&b);
```

Summary of pthreads

- POSIX Threads are based on OS features
 - can be used from multiple languages (need appropriate header)
 - familiar language for most of program
 - ability to shared data is convenient
- Pitfalls
 - data race bugs are hard to find because they can be intermittent
 - deadlocks are usually easier, but can also be intermittent
- OpenMP is commonly used today as an alternative